

SCOTT KOSKOSKI

Memphis, TN | scottkoskoski@gmail.com | 501-470-6677
[linkedin.com/in/scott-koskoski](https://www.linkedin.com/in/scott-koskoski) | github.com/scottkoskoski

SUMMARY

Machine Learning Engineer / Data Scientist with **6+ years** building production-scale ML, NLP, and optimization systems across retail, e-commerce, and pharma logistics. Deployed a GPU-accelerated freight platform (1M+ shipments; **multimillion-dollar savings**) and a FastAPI + FAISS RAG service (150k docs, <200ms latency). M.S. CS (in progress), M.A. Economics. Python expert. **Proven record of translating data into business impact.**

TECHNICAL SKILLS

| | |
|----------------------------|--|
| AI & NLP | Transformers (BERT, DeepSeek-V3), RAG, FAISS, LangChain, Hugging Face, BGE Embeddings |
| ML & Data | PyTorch (+CUDA), TensorFlow, scikit-learn, XGBoost, LightGBM, Pandas, NumPy, PySpark, Statsmodels |
| Languages & DBs | Python, SQL (T-SQL, PL/pgSQL), JavaScript/TypeScript, C/C++, R, SQL Server, PostgreSQL, Oracle, SQLite |
| Cloud & DevOps | Azure (VMs, Functions, Blob, ML), AWS (S3, EC2, Lambda), Docker, Git, GitHub Actions |
| Web & Viz | Flask, FastAPI, React, Next.js, REST/JSON, JWT Auth, Tableau, Matplotlib, ggplot2 |

PROFESSIONAL EXPERIENCE

GREEN MOUNTAIN TECHNOLOGY · Memphis, TN (Remote)

Senior Solutions Engineer (Data Scientist) | Apr 2025 – Present

- **Architected** a Cost Efficiency Index (CEI) platform ingesting LTL freight data from 6+ SQL Server DBs and training 5 GPU-accelerated models on 1M+ records with automated hyperparameter tuning.
- **Delivered** a minutes-to-insight workflow adopted by 6 enterprise shippers, replacing multi-day analyses with real-time freight-cost optimization.

Solutions Engineer II (Data Scientist) | Oct 2022 – Mar 2025

- **Engineered** Python anomaly-detection, cutting manual monitoring 10+ hrs/week and boosting data reliability 35%.
- **Led** RFP analyses and optimized ETL pipelines (-40% runtime), enabling negotiations that reduced client spend 15% annually.

Solutions Engineer I (Data Analyst) | Aug 2020 – Apr 2022

- **Built** Tableau supply-chain dashboards on complex SQL for Fortune 500 shippers, elevating operational insight.
- **Automated** freight-rating scripts in Python, cutting workload 25% and eliminating manual errors.

AUTOZONE · Memphis, TN (Hybrid)

Data Analyst | May 2022 – Oct 2022

- **Optimized** product-assortment algorithms via statistical analysis (Python, SAS, SQL), improving inventory efficiency 12%.
- **Designed** A/B-test frameworks across 5k+ stores and built causal & multivariate models to surface performance drivers.

VERISK FINANCIAL ARGUS (now part of TransUnion) · White Plains, NY

Analytics Associate – Merchant Analytics | Jun 2019 – Jul 2020

- **Forecasted** revenue with ARIMA & GAM models (92% accuracy) and deployed ML pipelines in R & Python.
- **Tuned** SQL queries across 15+ PB of transaction data and built interactive Tableau dashboards, delivering real-time insights.

EDUCATION

M.S. Computer Science — University of Tennessee, Knoxville | *Anticipated May 2026*

M.A. Economics — University of Memphis | *May 2019*

- **Outstanding** Economics Graduate Student 2019 · **Winner**, Wells Fargo Analytics Challenge

B.A. Economics, Minor Mathematics — University of Arkansas, Fayetteville | *May 2017*

SELECT PROJECT HIGHLIGHTS

- **RAG-Lab** – End-to-end FastAPI service answering U.S. employment questions using FAISS vector search (BGE-large embeddings, 150k records) and DeepSeek-V3 LLM; delivers cited responses with <200ms vector lookup.
- **DDoS Detection Pipeline** – Open-source PySpark pipeline ingesting PCAP traffic data, engineering features, and training classifiers to identify DDoS attacks; containerized with Docker.
- **RISC-V RV32IM Virtual Machine** – Educational VM implementation with 32 registers, 1 MiB memory, ELF loader, and complete fetch-decode-execute cycle in C.